


A *Markov Decision Process* (MDP) model contains:

- A set of possible world states S
- A set of possible actions A
- A real valued reward function $R(s,a)$
- A description T of each action's effects in each state.



We assume the *Markov Property*: *the effects of an action taken in a state depend only on that state and not on the prior history.*